

Data-driven Intelligent Systems

Lecture 11 Theory of Learning, Evaluation




<http://www.informatik.uni-hamburg.de/WTM/>

Theory of Learning from Data

- ▶ Model Learning
 - Statistical Learning Theory (VC Dimension, ERM, SRM)
 - Cost Function and its Bayesian View
- Training, Validation & Test Data
 - Cross Validation
- Evaluation of Classification models
 - Confusion Matrix

Machine Learning & Human Learning

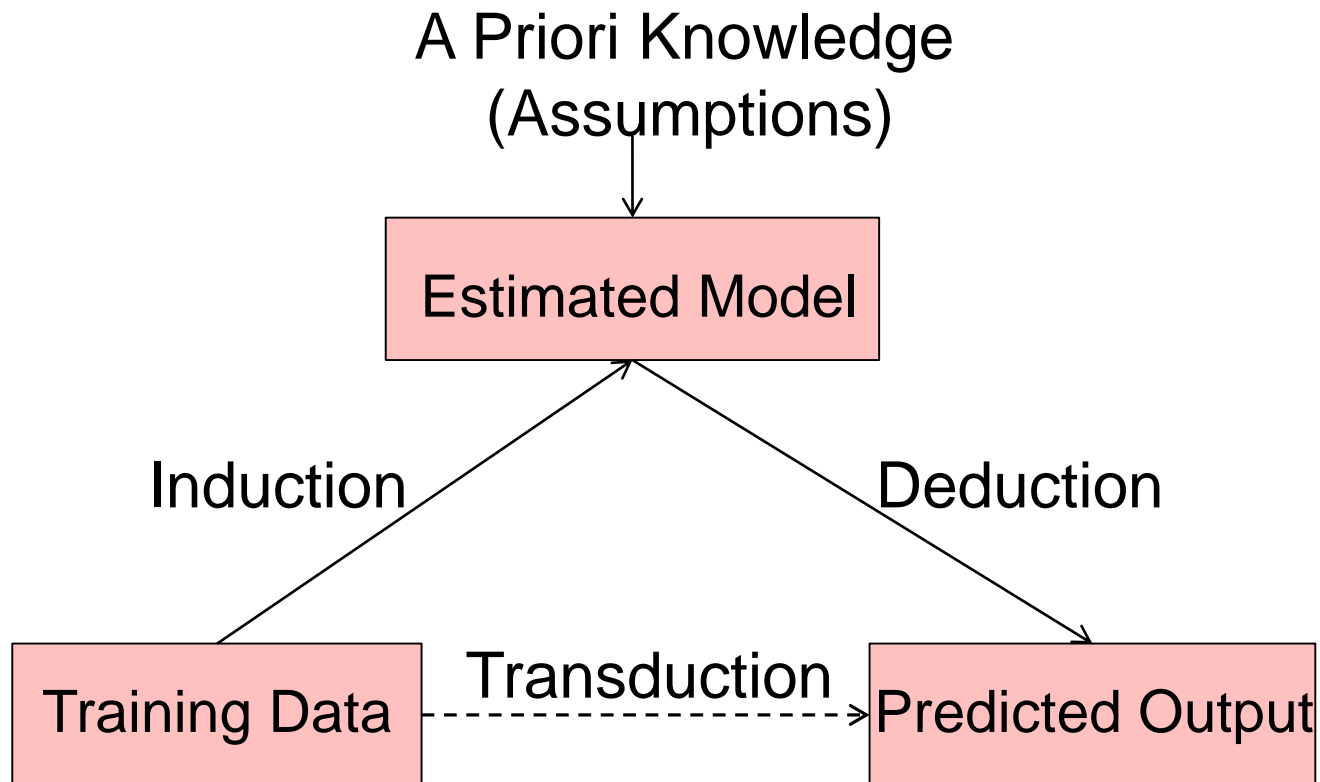
- Supervised, unsupervised, semi-supervised, self-supervised, reinforcement learning
- Learning from examples
- Case-based learning
- Learning by analogy
- Learning by doing
- ...



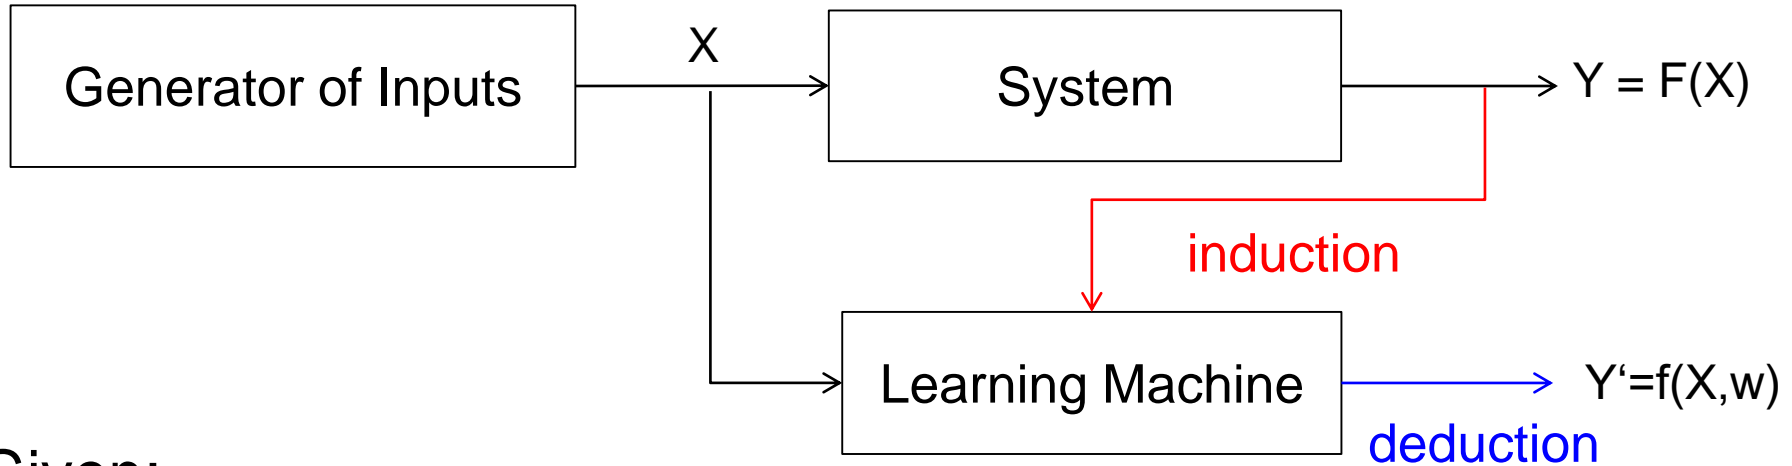
Machine Learning Issues

- Static vs. dynamic data
- Centralized vs. distributed data
- Batch vs. incremental (on-line) learning
- Active/adaptive learning
- Life-long learning
- ...

Types of Inference: Induction, Deduction, Transduction



A Learning Scenario



Given:

- observed samples $\{(X, Y)\}$

How to select $f(X, w)$:

- Approximating function f ?
 - Hyperparameters?
- Parameters: w ?

← A priori knowledge required!

Example:

f : linear in parameters:


$$y = w_1 x^n + w_2 x^{n-1} + \dots + w_0$$

nonlinear in parameters:




$$y = e^{-wx}$$

Hypotheses for a Given Data Set

- Given: samples (x_d, y_d)
- Unknown: true function $y=F(x)$
- Wanted: approximation $\underbrace{h(x)}_{\text{hypothesis}}$ of the true function



Polynomial (linear, quadratic, etc.) or exponential model?



How to Learn with a Learning Machine? (1)

- Learning objective
 - *Inductive principle* – a general prescription for learning
 - Tells us **what** we wish to achieve with the data
 - define a Risk function
 - choose a model (approximating function) of suitable complexity
- Learning method
 - Tell us **how** to obtain an optimal estimate
 - I.e. a constructive implementation of an inductive principle
 - find good model parameters

How to Learn with a Learning Machine? (2)

- **Loss function** $L(y_d, f(x_d, w))$: (also: **Error function**)
 - measure of difference between y_d and $f(x_d, w)$ for each sample d
 - y_d – the output produced by the system,
 - X_d – a tuple of inputs,
 - $f(X_d, w)$ – the output produced by the learning machine for a selected approximating function f ,
 - w – the set of parameters in the approximating function.
- **Risk functional** $R(w)$:
 - measure of accuracy of the learning machine:

$$R(w) = \frac{1}{\#d} \cdot \sum_d L(y_d, f(x_d, w))$$

Analogue terms: **Cost**, Score, Profit, Fitness,
Utility, Reward, **Objective** function

How to Learn with a Learning Machine? (3)

- Examples of loss function $L(y, f(x, w))$:

- **Classification error:**

$$L(y, f(x, w)) = \begin{cases} 0 & \text{if } y = f(x, w) \\ 1 & \text{if } y \neq f(x, w) \end{cases}$$

- **Squared error** (a measure for regression):

$$L(y, f(x, w)) = (y - f(x, w))^2$$

Theory of Learning from Data

- Model Learning
 - ▶ Statistical Learning Theory (VC Dimension, ERM, SRM)
 - Cost Function and its Bayesian View
- Training, Validation & Test Data
 - Cross Validation
- Evaluation of Classification models
 - Confusion Matrix

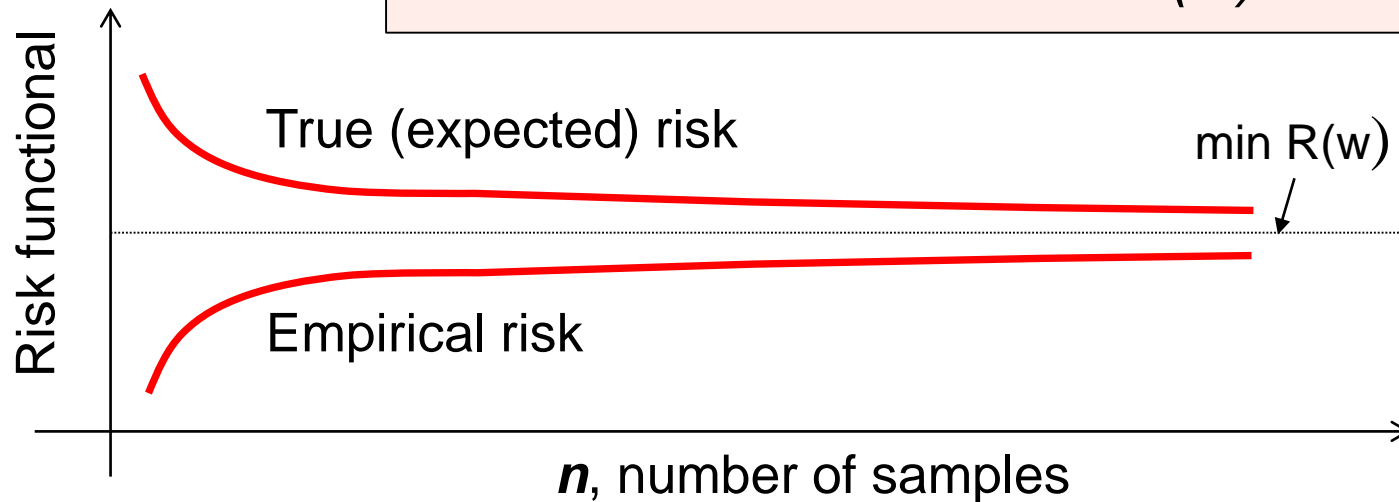
Statistical Learning Theory (1)

- SLT – formalizes many learning procedures developed in AI, ANN, statistics, Data Mining, Pattern Recognition
- SLT considers learning with small sets of samples
 - Exact distribution of data $p(x, y)$ is unknown
 - When does overfitting occur?
 - Approximate true risk $R(w)$ with an empirical risk
- **Empirical Risk Minimization (ERM)** – the basic inductive principle:
 - Find the optimal estimate = minimum of risk function $R(w)$ based only on the available data
 - Implementation of ERM depends on selected L and $f(x, w)$
- SLT = VC theory (**V**apnik **C**hervonenkis)

Statistical Learning Theory (2)

- Asymptotically consistent estimator:

when $n \rightarrow \infty$ then true $R(w) \approx$ empirical risk



Assume:


- given data distribution
- fixed number of model parameters
- model fully trained for each n

- For $n \rightarrow \infty$
 - true model parameter values will be estimated
 - model will generalize to “unseen” data
- Asymptotic consistency should hold for ALL classes of approximating functions

Statistical Learning Theory (3)

- To ensure ***asymptotic consistency***, approximating functions should be like a ***growth function***
- As the number of samples grows, the approximating functions should start to ***generalize***
- Generalization means
 - failure to model noise
 - failure to model overly complex data
- The ***set of hypothesis*** that the approximating function can make over the data ***should be limited***

Growth Function



- Hypothesis set H = all the functions a learner can approximate
- A growth function is defined as


$$\Pi_H(n) = \max |H(S)|$$

over all input sets S of size n

i.e. the maximum
number of ways
 n points can be
classified by H

- E.g. binary classification: $\Pi_H(n) \leq 2^n$

Vapnik Chervonenkis (VC) Dimension




VC dimension: Point $n = VC$ where growth starts to slow down

- The VC dimension of H is the cardinality of the largest set S that can be fully represented by H (i.e. learned)
- $VC(H)$ is typically finite in good learners
- A “saturating” growth function ensures asymptotic consistency

VC Dimension, Examples

- Linear classifier in 2D:

$$VC(H) = 3$$



- Linear classifiers for d features plus a constant term b :

$$VC(H) = d+1 \quad (\text{perceptron})$$

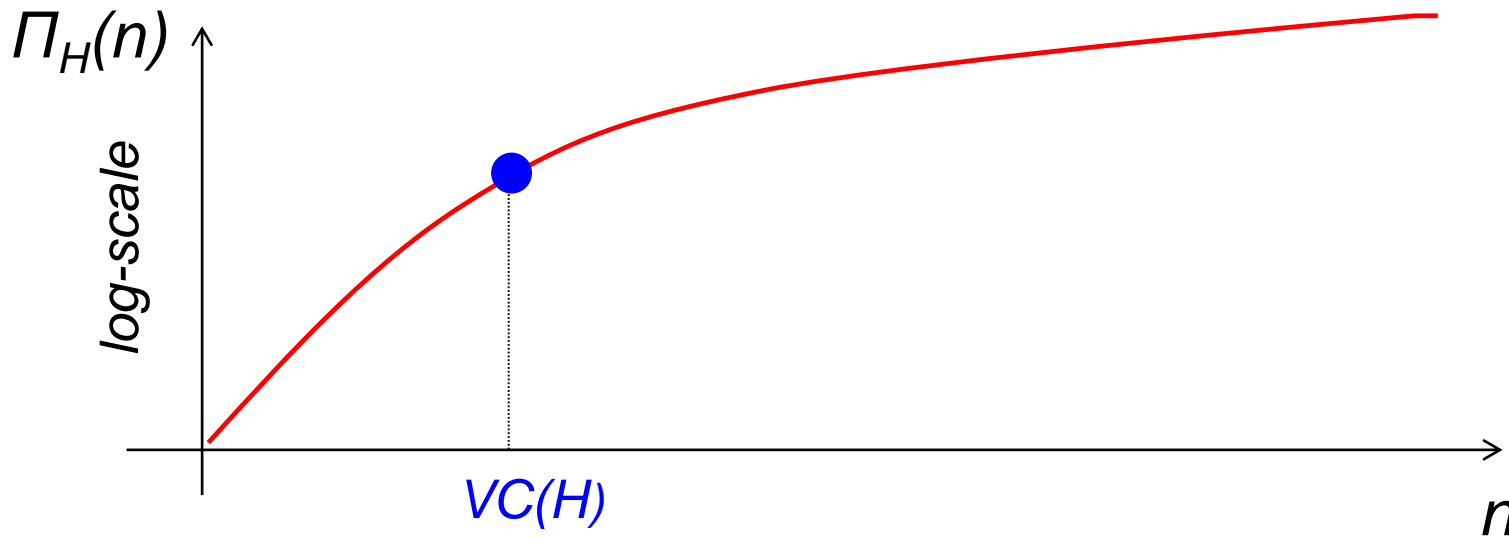
- Neural networks:

$$VC(H) \approx \# \text{parameters}$$

- Decision tree of rank r that defines Boolean functions on n boolean variables:

$$VC(H) = \sum_{i=0}^r \binom{n}{i}$$

VC Dimension




VC dimension: Point $n = VC$ where growth starts to slow down

- **ERM** applicable for large n ($n/VC > 20$)
- Possible overfitting for small n ($n/VC < 20$)
→ need to constrain the structure of the learner → **SRM**

Structural Risk Minimization (1)

- SRM requires **a priori** specification of a structure for sets of approximating functions S_1, S_2, \dots, S_k .



$$VC(S_1) < VC(S_2) < \dots < VC(S_k)$$

- **SRM approach** towards optimal model:
 - Calculate or estimate VC-dimension for any element S_k
 - Minimize empirical risk $R(w)$ for each S_k
- The optimal solution is a tradeoff:
 - High complexity (large VC)
→ small empirical risk
 - Low complexity (small VC)
→ empirical risk \sim true risk


good
genera-
lization

Structural Risk Minimization (2)

- SRM – a trade off between **complexity** (of approximating functions) and **quality** (of results)


”As simple as possible,
but with enough quality.“

Occam's razor principle



- Optimal model estimation:
 - Select an element of the structure with optimal complexity
 - Define the model based on selected approximating functions
 - Penalize complex models by **regularisation**


SRM Optimization Strategy



With increasing complexity of approximating functions true & empirical risk $R(w)$ decrease until the value – VC dimension; thereafter they diverge.

- Optimization:
 - Stochastic approximation (or gradient descent)
 - Iterative methods
 - Greedy optimization (following locally optimal choice)

Complexity and Generalization



- Complexity = degrees of freedom in the model
 - **E.g.:** number of variables
 - Effective model complexity may rise over the course of training (this justifies **early stopping**)

Bias-Variance Tradeoff

■ Model *bias*:

may result from SRM

- Model outputs are often biased – models can learn certain aspects of the data, but have limitations elsewhere
 - *Underfitting* is a form of bias
- Model bias unwanted, since output shall depend on the data
- But: a smart bias may enable certain model performance!

■ Model *variance*:

may result from ERM

- Models' outputs often have large variance under:
 - small variations in the data, e.g. different sampling, or
 - with different initial random values of the model parameters
- Unwanted variance often observed in powerful models, which are unconstrained by model bias
 - *Overfitting* models have this behaviour

Theory of Learning from Data

- Model Learning
 - Statistical Learning Theory (VC Dimension, ERM, SRM)
 - ▶ Cost Function and its Bayesian View
- Training, Validation & Test Data
 - Cross Validation
- Evaluation of Classification models
 - Confusion Matrix


Bayes

- Probability distribution of two random variables:

$$\begin{aligned} P(D, H) &= P(D | H) \cdot P(H) \\ &= P(H | D) \cdot P(D) \end{aligned}$$

- Rearrange terms:

$$P(H | D) = \frac{P(D | H) \cdot P(H)}{P(D)}$$



Relation of ERM/SRM to Bayesian View

Bayes Theorem:
$$P(H | D) = \frac{P(D | H) \cdot P(H)}{P(D)}$$

- $P(D|H)$: **Likelihood** that model H generates the data D .
Find maximum likelihood model ~ **ERM**
~ try to model the data best, at any price.
- $P(H)$: **Prior** probability of model H ; **penalizes models of complex structure**; based on *a priori* knowledge.
- $P(D)$: Evidence; just a normalizing factor.
- $P(H|D)$: **Posterior** probability for H after having seen the data.
Find maximum posterior model
Tradeoff: well performing & simple.
~ **SRM**

Relation Likelihood vs. Empirical Risk

- Likelihood for the model to generate the data:

maximise \longrightarrow
$$P(D | H) \sim \prod_{\text{data } d} e^{-(y_d - f(x_d, w))^2}$$

Gaussian prob. of data
deviating from model


- Take $-\ln(\cdot)$ on both sides of the equation:
- Formulation as cost function:

minimise \nearrow
$$-\ln(P(D | H)) \sim \sum_{\text{data } d} (y_d - f(x_d, w))^2$$

Square error
(empirical) Risk

Relation Likelihood vs. Empirical Risk

Likelihood probability
(of data being produced by the model)



Assumption:
Gaussian-deviated
data around model

error
(between data and prediction)

Probabilities vs. Cost Functions

- Bayes probabilistic formulation:

maximise \longrightarrow

$$\underset{\substack{\nearrow \\ \text{Posterior}}}{P(H | D)} \sim \underset{\substack{\nwarrow \\ \text{Likelihood}}}{P(D | H)} \cdot \underset{\substack{\nwarrow \\ \text{Prior}}}{P(H)}$$

- Take $-\ln(\cdot)$ on both sides of the equation:

minimise \longrightarrow

$$\underset{\substack{\nearrow \\ \text{Structural} \\ \text{Risk}}}{-\ln(P(H | D))} \sim \underset{\substack{\nwarrow \\ \text{Empirical} \\ \text{Risk}}}{-\ln(P(D | H))} - \underset{\substack{\nwarrow \\ \text{Penalty on parameters} \\ \text{(regularizer)}}}{\ln(P(H))}$$

\longrightarrow *Maximising the posterior probability* of the model is equivalent to *minimising costs*.

Probabilities vs. Cost Functions – Example

- Probabilistic formulation:

maximise \longrightarrow
$$P(H | D) \sim \prod_{\text{data } d} e^{-(y_d - f(x_d, w))^2} \cdot e^{-w^2}$$

\nwarrow Gaussian prob. of data deviating from model

\nwarrow Small parameters w have larger prior probability

- Formulation as cost function:

minimise \nearrow
$$-\ln(P(H | D)) \sim \sum_{\text{data } d} (y_d - f(x_d, w))^2 + w^2$$

\nearrow Square error

\nwarrow Penalty on large w

imposes a model bias

Theory of Learning from Data

- Model Learning
 - Statistical Learning Theory (VC Dimension, ERM, SRM)
 - Cost Function and its Bayesian View
- ▶ Training, Validation & Test Data
 - Cross Validation
- Evaluation of Classification models
 - Confusion Matrix

Using Data

- Use this data to **find the best parameters** w for each model k


$$f_k(x, w)$$

- Use this data to calculate an estimate of **score** $S_k(w)$ for each $f_k(x, w)$ **and select**

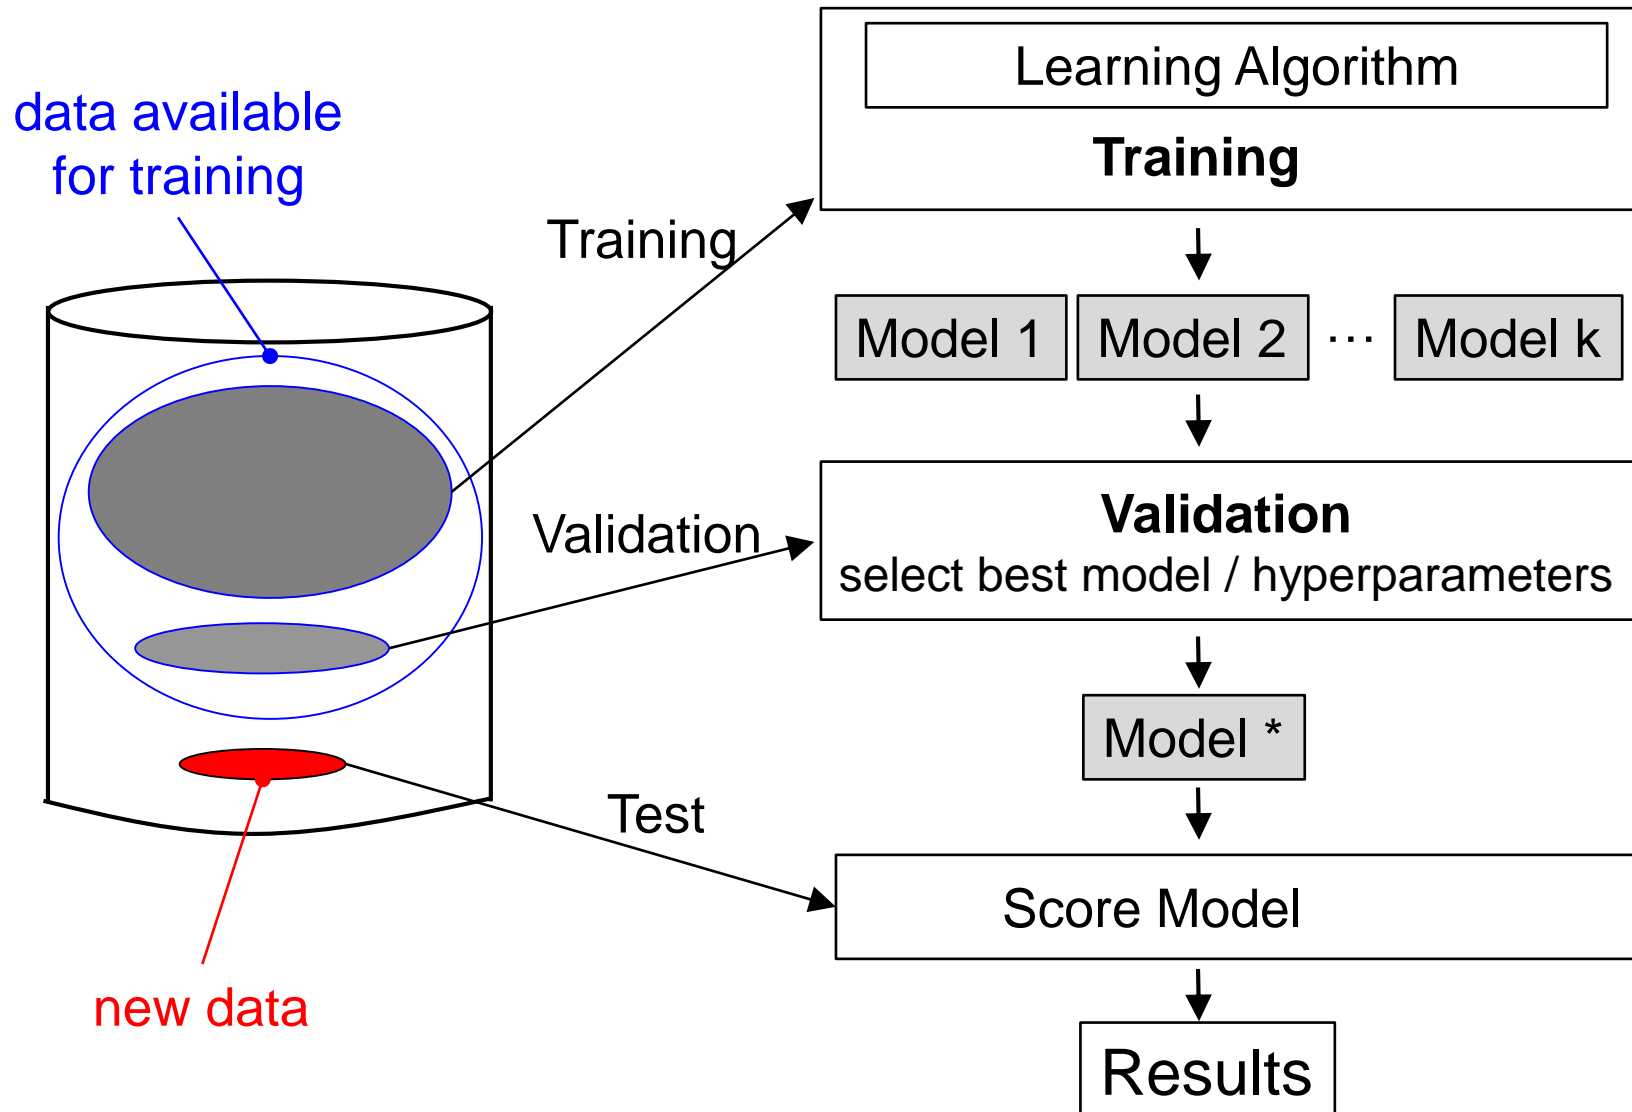
$$k^* = \operatorname{argmin}_k S_k(w)$$

→ **find best hyperparameters**

- Use this data to calculate an **unbiased estimate** of $S_{k^*}(w)$ for the selected model



The Data Mining Process



Theory of Learning from Data


- Model Learning
 - Statistical Learning Theory (VC Dimension, ERM, SRM)
 - Cost Function and its Bayesian View
- Training, Validation & Test Data
 - ▶ Cross Validation
- Evaluation of Classification models
 - Confusion Matrix

Making Most of the Data

- **Resubstitution** method:
 - training data = testing data; naïve strategy, optimistically biased; not for small n .
- **Bootstrap** method:
 - resample randomly with replacement to generate data sets of same size but different proportion of samples for training and testing.
- **Holdout** method:
 - $x\%$ of data for training , $(1-x)\%$ for testing.
- **Rotation** method (**k -fold cross validation**):
 - total of k data segments, $k-1$ for training, one for testing; repeat k times.
- **Leave-one-out** method:
 - $n-1$ training samples, one testing sample; repeat n times.

Hold-out Method

- **Hold-out set.** Partition data into training and test sets



- Data from the test set are “lost” for training
- Different partitioning → different estimates

K-fold Cross Validation

- Create K equal partitions

- **Example 1:**



10-fold cross validation:

- Use the first 90% of the data set for training and then test on the final 10%
- Then use the next 10% for testing etc.


- **Example 2:**

$K=n$, number of data points

- “Leave-one-out method”
- Train n-times with n-1 data points



K-fold Cross Validation (for Regression)



Linear Regression:
 $MSE_{3FOLD} = 2.05$


Randomly break the dataset into
k partitions
(here: k=3 – red, blue, purple)

- For red=test: Train on the points **not** in the red partition. Find the test-sum of errors on the red points.
- For blue=test: Train on the points **not** in the blue partition. Find the test-sum of errors on the blue points.
- For purple=test: Train on the points **not** in the purple partition. Find the test-sum of errors on the purple points.

Then report the
mean square
error (MSE).

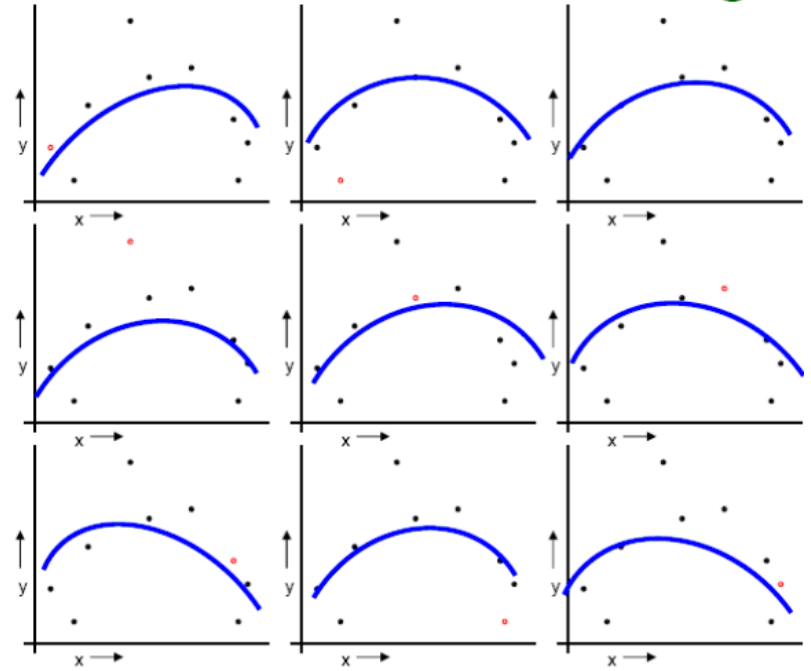
Examples: Leave One Out Cross Validation

Linear regression (2 parameters)



$$\text{MSE}_{\text{LOOCV}} = \mathbf{2.12}$$

Quadratic regression (3 parameters)



$$\text{MSE}_{\text{LOOCV}} = \mathbf{0.962}$$

→ quadratic model is better: better hyperparameters

MSE typical for regression.

Which measure for classification?


Theory of Learning from Data

- Model Learning
 - Statistical Learning Theory (VC Dimension, ERM, SRM)
 - Cost Function and its Bayesian View
- Training, Validation & Test Data
 - Cross Validation
- ▶ Evaluation of Classification models
 - Confusion Matrix


Evaluation of Classification Systems (1)

- Task: Determine which of a fixed set of classes an example belongs to.
- Input: Training set of examples annotated with class values.
- Output: Induced hypothesis (model/concept description/classifier).

Learning



Classification




Evaluation of Classification Systems (2)

Evaluation criteria:

- ***Accuracy*** of the classification
- ***Interpretability***
 - E.g. size of a decision tree; insight gained by the user
- ***Efficiency***
 - ... of model construction
 - ... of model application
- ***Scalability***
 - ... for large datasets
- ***Robustness***
 - w.r.t. noise and unknown attribute values

Evaluation of Classification Systems (3)

- Training set: examples with class values for learning.
- Test set: examples with class values for evaluating.
- **Evaluation**: Model hypotheses are used to classify the test data; results are compared to known classes.



- **Accuracy**: percentage of examples in the test set that is classified correctly.
- Binary classification: “positive” or “negative”

Theory of Learning from Data

- Model Learning
 - Statistical Learning Theory (VC Dimension, ERM, SRM)
 - Cost Function and its Bayesian View
- Training, Validation & Test Data
 - Cross Validation
- Evaluation of Classification models
 - ▶ Confusion Matrix

Classifier Evaluation Metrics: Accuracy & Error Rate

- **Confusion Matrix:**

Predicted class\Actual class	C_1	$\neg C_1$
C_1	True Positives (TP)	False Positives (FP)
$\neg C_1$	False Negatives (FN)	True Negatives (TN)

- **Classifier Accuracy**, or recognition rate: percentage of test set tuples that are correctly classified,

$$\text{accuracy } A = \frac{TP + TN}{TP + TN + FP + FN}$$

- **Error rate:** $1 - \text{accuracy}$, or

$$\text{error rate} = \frac{FP + FN}{TP + TN + FP + FN}$$

Classifier Evaluation Metrics

- **Sensitivity/Recall:** True Positive recognition rate

=1 if *all data*
classified
as positive

$$R = \frac{TP}{TP + FN} \quad (TP + FN = \text{actual positives})$$

- **Specificity:** True Negative recognition rate

$$SP = \frac{TN}{TN + FP} \quad (TN + FP = \text{actual negatives})$$

- **Precision:** exactness – what % of tuples that the classifier labelled as positive are actually positive?

$$P = \frac{TP}{TP + FP}$$

=1 if *just one* data
point safely classified
as positive

- Perfect score is 1.0

- **Opposing goals** when maximising **precision**
& **recall**

Classifier Evaluation Metrics:

F Measure

- ***F* measure** (F_1 or ***F*-score**): harmonic mean of precision and recall

$$F = \frac{2 \cdot \textit{precision} \cdot \textit{recall}}{\textit{precision} + \textit{recall}}$$

Confusion Matrix / Metrics – Summary

Actual Predicted	Class 1	Class 2
	Class 1	Class 2
Class 1	True Positive	False Positive
Class 2	False Negative	True Negative

Evaluation metrics:

Accuracy

$$A = (TP + TN) / (TP + FP + FN + TN)$$

TP rate, Sensitivity, **Recall**

$$R = TP / (TP + FN)$$

FP rate

$$FPr = FP / (FP + TN) = 1 - \text{TN rate}$$

TN rate, Specificity

$$SP = TN / (FP + TN) = 1 - FPr$$

Precision

$$P = TP / (TP + FP)$$

F-score

$$F = 2 P \cdot R / (P + R)$$

Classifier Evaluation Metrics:

Example Confusion Matrix

Actual class\ Predicted class	buy_computer = yes	buy_computer = no	Total	Recognition (%)
buy_computer = yes	6954	412	7366	99.34 <i>sensitivity</i>
buy_computer = no	46	2588	2634	86.27 <i>specificity</i>
Total	7000	3000	10000	95.42 <i>accuracy</i>

- Given m classes, an entry, $\mathbf{CM}_{i,j}$ in a **confusion matrix** indicates # of tuples in class i that were labeled by the classifier as class j .
- Extra rows/columns may provide totals or recognition rate per class.

Confusion Matrix for Three Classes

Classification Model	True Class			Total
	0	1	2	
0	28	1	4	33
1	2	28	2	32
2	0	1	24	25
Total	30	30	30	90

$$\text{Error} = \frac{\text{Sum of non diagonal}}{\text{Total}} = 10 / 90 = 0.11 \text{ (11\%)}$$

$$\text{Accuracy} = 1 - \text{Error} = 1 - 0.11 = 0.89 \text{ (89\%)}$$

Accuracy Unsuitable for Skewed Distributions

- Typical *Class Imbalance Problem*: majority in negative class

P\A	C1	C2
C1	0	0
C2	10	90



Accuracy	90/100
Recall (sensitivity)	0/10
Precision	0/0
F-Score	0/0

P\A	C1	C2
C1	3	10
C2	7	80



Accuracy	83/100
Recall	3/10
Precision	$\frac{3}{13}$ =0.23
F-Score	$\frac{6}{23}$ =0.261

P\A	C1	C2
C1	8	42
C2	2	48



Accuracy	56/100
Recall	8/10
Precision	$\frac{8}{50}$ =0.16
F-Score	$\frac{4}{15}$ =0.267

Cost Matrix

	ACTUAL CLASS		
	$c(i j)$	Class=Yes	Class=No
	Class=Yes	C(Yes Yes)	C(No Yes)
	Class=No	C(Yes No)	C(No No)

$c(i | j) = c_{ij}$ – Cost of misclassifying class j example as class i

Total cost function:
$$C = \sum_i \sum_j c_{ij} \cdot e_{ij}$$

Computing Cost of Classification

<i>Cost Matrix</i>	Actual Class	
Predicted Class	C(i j)	
		+ -
+	-1	1
-	100	0

<i>Model M₁</i>	Actual Class	
Predicted Class		+ -
+	150	60
-	40	250

Accuracy = 80%
Cost = 3910

<i>Model M₂</i>	Actual Class	
Predicted Class		+ -
+	180	160
-	10	150

Accuracy = 66%
Cost = 980

Summary

- Statistical learning theory provides a theoretical foundation why a more powerful model isn't always better
 - Parallels between probabilistic (Bayes) and cost function formulations
- Validation and test data may come costly
 - cross validation makes the most of available data
- Confusion matrix and various evaluation metrics
 - accuracy, precision, recall, F-score