

Data-driven Intelligent Systems

Lecture 14 Ensemble Learning 1



<http://www.informatik.uni-hamburg.de/WTM/>

Ensemble Learning – Overview



Benefits of ensembles

- How to combine their outputs
- Bagging
- Boosting
 - AdaBoost

Ensemble Learning

- So far – learning methods learn a single hypothesis (model), chosen from a hypothesis space to make predictions
- ***“There ain’t no such thing as a free lunch”***
 - No single algorithm wins all the time!
- Ensemble learning
 - select a collection (*ensemble*) of hypotheses (*models*) and combine their predictions
- **Example:** Generate 100 different decision trees from the same or different training set and have them vote on the best classification for a new example.



Photo © Jon Worth / atheistbus.org.uk

Value of Ensembles

- Key motivation: **reduce** the **error rate**!
Hope: it is less likely that an *ensemble* misclassifies an example
- **Examples**: Human ensembles are demonstrably better:
 - How many jelly beans in the jar?:
Individual estimates vs. group average
 - Who Wants to be a Millionaire: Audience vote
 - Diagnosis based on multiple doctors' majority vote
- **Theory behind**: We combine multiple **independent** and **diverse** decisions
 - each is at least more accurate than random guessing
 - random errors cancel each other out
 - correct decisions are more consistent and add up

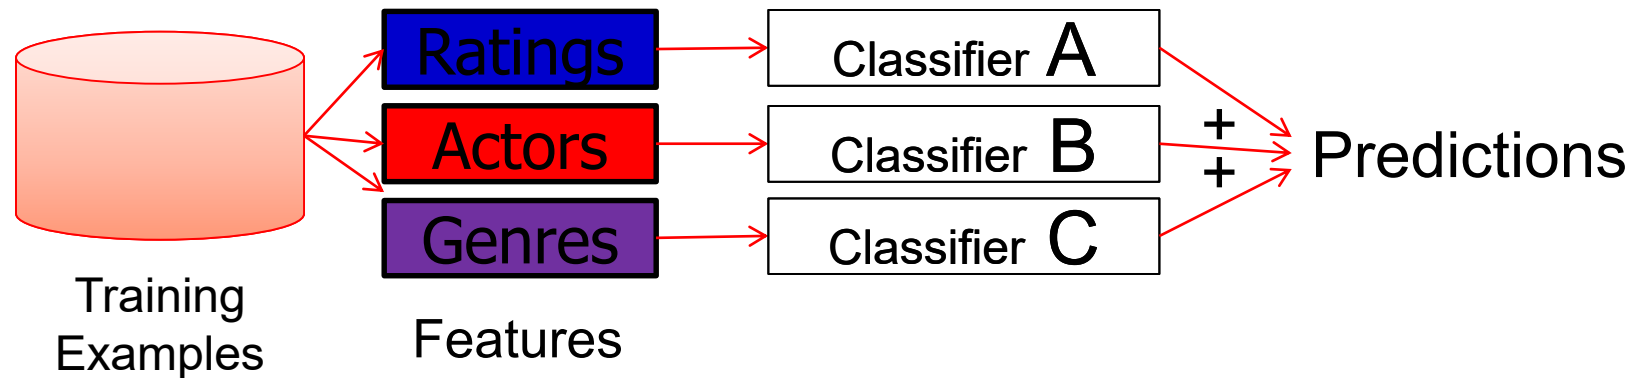


Achieving Diversity (1)

1. Using different learning **algorithms**
← *how many algorithms do we know?*
2. Using different **hyper-parameters** in the same algorithm
← *some parameters not as good as others*
3. Using different **input representations**, e.g. different subsets of input *features*
← *sometimes: diversity hand-designed*
← *requires redundant features*
(e.g.: **Random Subspace Method**)
4. Using different **subsets of training data**
 - e.g. **bagging**, **boosting**, and **cascading**
← *diversity achieved automatically*

Achieving Diversity (2)

3. Diversity from *differences in input features*:



Example: multimodal emotion classification:

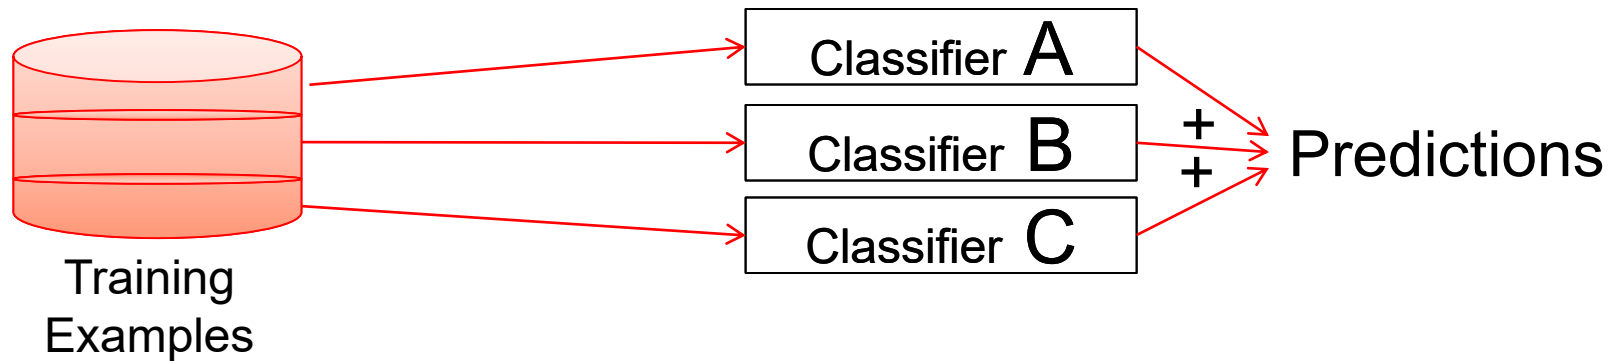
Model 1 ← vision; Model 2 ← audio; Model 3 ← text

Automatic feature selection:

Random Subspace Method, works for large feature sets with redundant features

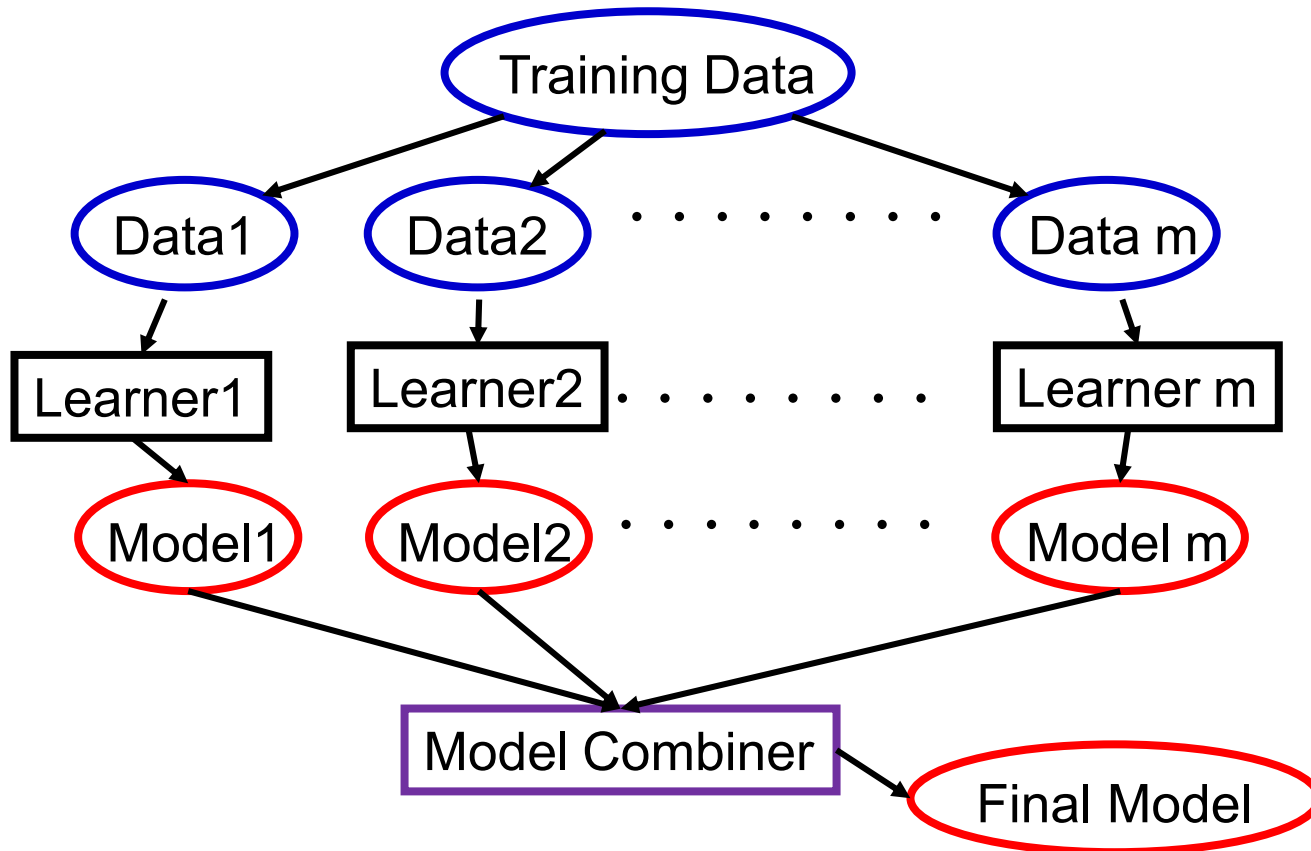
Achieving Diversity (3)

4. Diversity from *subsets of training data*:
use ***different subsets of training data to*** learn multiple alternative definitions of a concept



Achieving Diversity (4)

4. Diversity from *subsets of training data*:
use *different subsets of training data to* learn multiple alternative definitions of a concept



Ensemble Learning – Overview

- Benefits of ensembles
- ▶ How to combine their outputs
- Bagging
- Boosting
 - AdaBoost

How to Combine the Outputs of Base Learners?

- **Global approach** is through fusion – the outputs of all learners are combined by *averaging*, *voting*, or *stacking**
- **Local approach** is based on *learner selection* – it examines the input and chooses the learner(s) responsible for generating the output
- **Multistage combination** use a serial approach where the next learner is trained with or tested on instances only where previous learners failed, or were inaccurate

**stacking*: a (simple) model
classifies the learners' *outputs*

Global Approach: Averaging Example

- Guess: how many people are in the picture?



- Is the ensemble average better than the individual estimates?

Averaging Example



How many? True number = T ; estimates = $\{x_i\}$, $i = 1, \dots, N$

- Individual errors: $e_i^{ind} = |T - x_i|$
 - averaged: $L1^{ind} = \frac{1}{N} \sum_i^N e_i^{ind}$
- Ensemble estimate: $m^{arith} = \frac{1}{N} \sum_i^N x_i$
 - Ensemble error: $L1^{ens} = |T - m^{arith}|$

$$sq_i^{ind} = (T - x_i)^2$$


















































$$SQ^{ind} = \frac{1}{N} \sum_i^N sq_i^{ind}$$

$$SQ^{ens} = (T - m^{arith})^2$$

We find: $L1^{ens} \leq L1^{ind}$ (equality if all $x_i \leq T$ or all $x_i \geq T$)

We find: $SQ^{ens} \leq SQ^{ind}$ (equality if all x_i equal)

Voting Example: Weather Forecast

Reality								
Learner's predictions	1							
	2							
	3							
	4							
	5							
Combine								

- Combine decisions of multiple models using *voting* procedure!

Voting: Lower Error Than Individuals

- Assume, binary base classifiers with error rate $\varepsilon = 0.3$
- Then, majority vote of $n=15$ independent classifiers:

$$\varepsilon_{ensemble} = \sum_{k=\text{ceil}(15/2)}^{15} \underbrace{\binom{15}{k}}_{\text{number of } k\text{-subsets in 15 items (binomial coefficient)}} \cdot \underbrace{\varepsilon^k (1 - \varepsilon)^{15-k}}_{\substack{\text{probability of } k \text{ outcomes in 15 draws, if } p(\text{outcome}) = \varepsilon \\ \rightarrow \text{probability that exactly } k \text{ classifiers make an error}}} = 0.05$$


$= 8$ sum over where more than half of the classifiers are wrong

- Binomial probability formula

Voting: Ensembles Give Better Results

- Majority vote of $n=15$ classifiers, error rate each $\varepsilon=0.3$:


$$\varepsilon_{ensemble} = \sum_{k=8}^{15} \underbrace{\binom{15}{k} \cdot \varepsilon^k (1 - \varepsilon)^{15-k}}_{\text{PMF, probability mass function}} = 0.05$$




Voting: Ensembles Give Better Results

- Majority vote of $n=15$ classifiers, error rate each $\epsilon=0.3$:

$$\epsilon_{ensemble} = \sum_{k=8}^{15} \binom{15}{k} \cdot \epsilon^k (1 - \epsilon)^{15-k} = 0.05$$



(a) Identical predictive models vs. different predictive models in an ensemble



(b) The different number of predictive models in an ensemble

Rank-Level Fusion Method

- Four-class problem (a,b,c,d)?

Rank / score	Classifier 1	Classifier 2	Classifier 3
4	c	a	d
3	b	b	b
2	d	d	c
1	a	c	a

$$r_a = r_a(C1) + r_a(C2) + r_a(C3) = 1 + 4 + 1 = 6$$

$$r_b = r_b(C1) + r_b(C2) + r_b(C3) = 3 + 3 + 3 = 9$$

$$r_c = r_c(C1) + r_c(C2) + r_c(C3) = 4 + 1 + 2 = 7$$

$$r_d = r_d(C1) + r_d(C2) + r_d(C3) = 2 + 2 + 4 = 8$$

- The winner-class is **b** because it has the maximum overall score

Global Approach: Combination Methods

Alternatives for combination are:

- Simple average (equal weights); for regression
- Majority voting; for classification
- Rank-level fusion; for multiple classes
- Average using weighted sum (unconstrained weights)
- Voting: linear combination of outputs d_j for learners j :

$$y = \sum w_j \cdot d_j \quad \text{where weights } w_j \geq 0 \quad \text{and} \quad \sum w_j = 1$$

with trained weights, this is an example of [stacking](#)

- Median
- Maximum or minimum
- Geometric mean: $\sqrt[k]{d_1 \cdot d_2 \cdot \dots \cdot d_k}$

Local Approach: Dynamic Classifier Selection

Algorithm

1. Define the local region
 - E.g. find the k nearest training points to the test input
2. Estimate the competence of the classifiers there
 - E.g. obtain accuracies of the classifiers on these points
3. Select
 - Choose the one that performs best on them
(or vote over a few “competent” ones).

Ensemble Learning – Overview

- Benefits of ensembles
- How to combine their outputs

Bagging

- Boosting
 - AdaBoost

Homogenous Ensembles

- Use a single, arbitrary learning algorithm but **manipulate training data** to make it learn multiple (diverse) models.
 - $\text{Data1} \neq \text{Data2} \neq \dots \neq \text{Data } n$
 - $\text{Learner1} = \text{Learner2} = \dots = \text{Learner } n$

Methods to change training data:

- **Bagging:**
 - Resample training data
- **Boosting:**
 - Reweight training data

Bagging: Bootstrap Aggregation (1)

- Training
 - Given a set D of tuples
 - At each iteration i , a **random subset** D_i of tuples is sampled *with replacement* from D for training (**bootstrap**) *
 - A classifier model M_i is learned for each training set D_i
- **Classification**: classify an unknown sample X
 - Each classifier M_i returns its class prediction
 - The bagged classifier M^* counts the votes and **assigns X to the class with the most votes**
- **Regression** (predict continuous outputs):
by taking the **average value** of each prediction for a given test sample
 - → *each set D_i is expected to have $\sim 2/3$ unique tuples and $\sim 1/3$ duplicates*

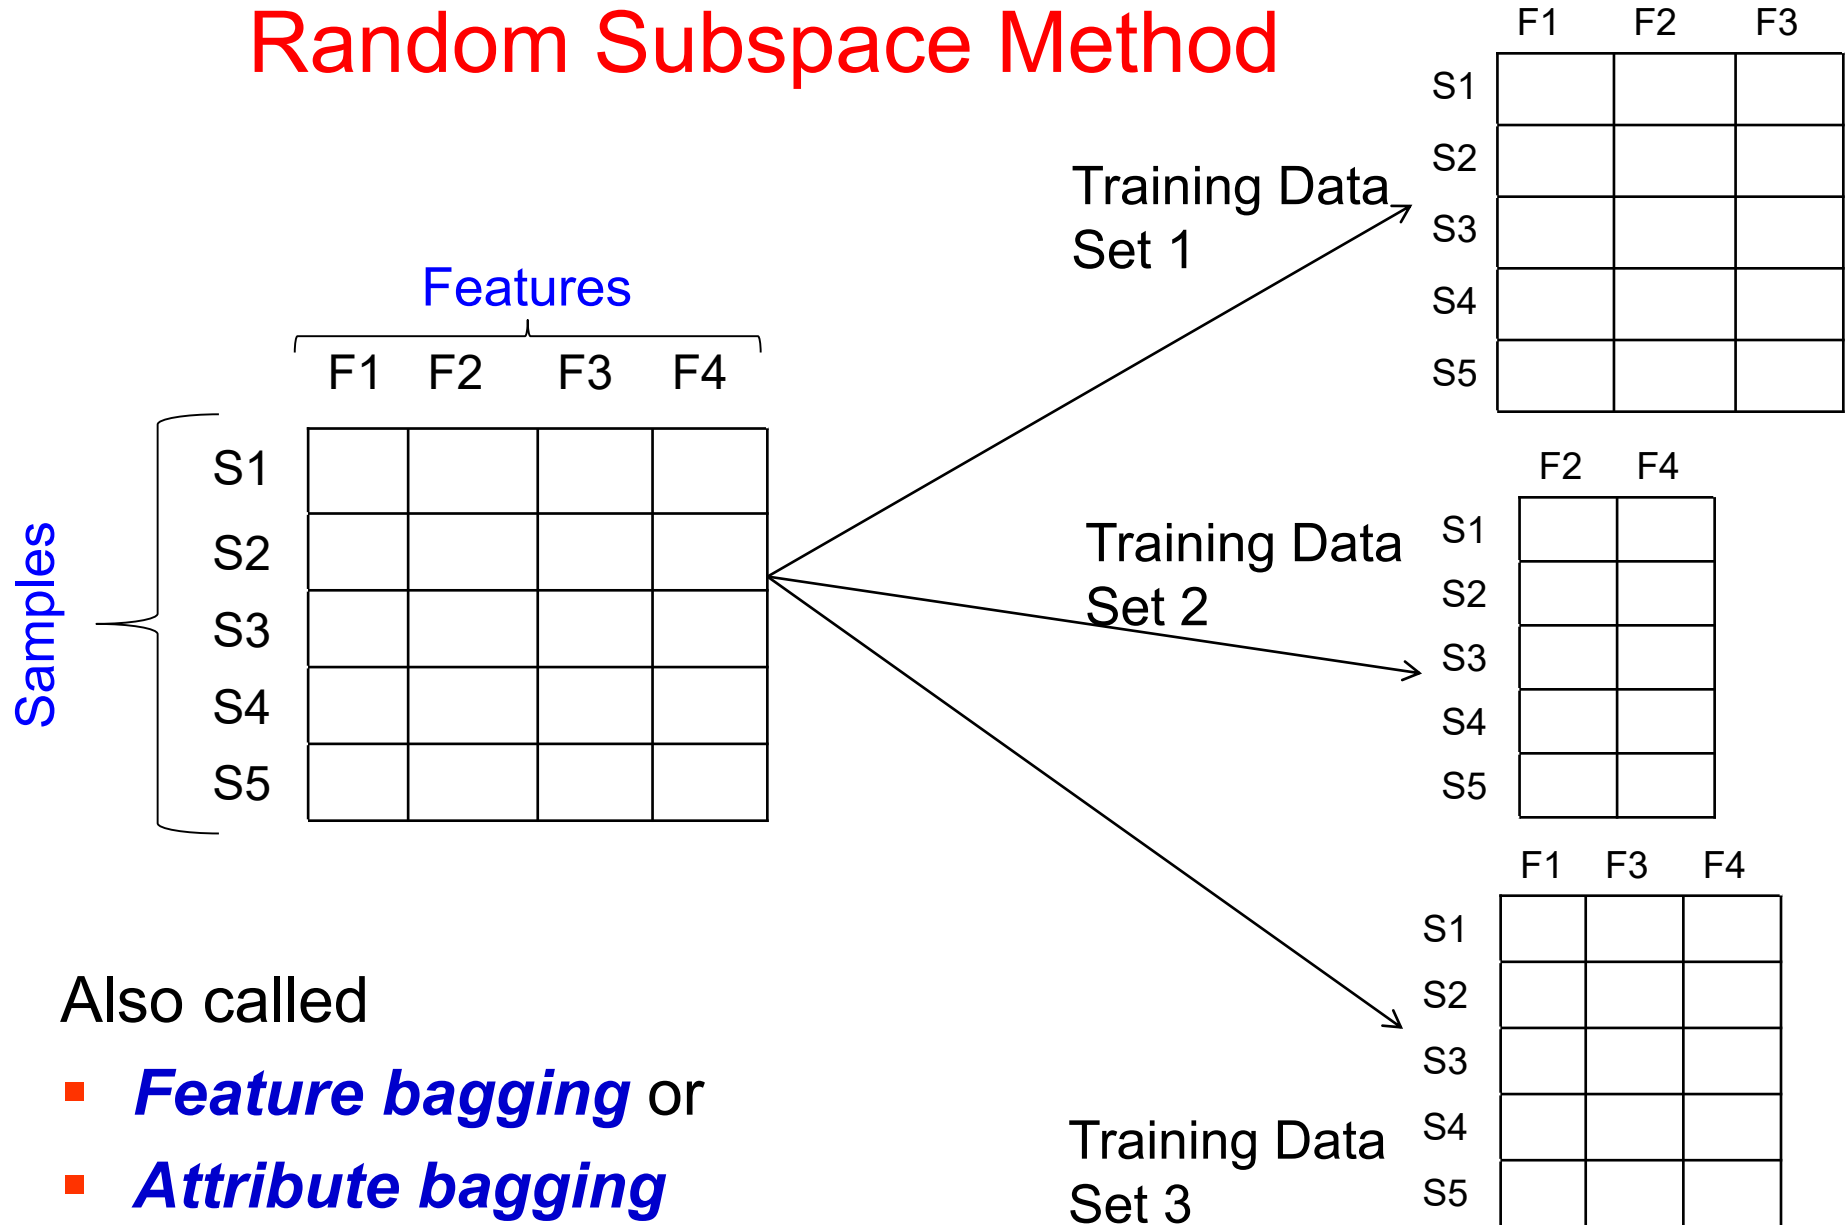
Bagging: Bootstrap Aggregation (2)

■ Accuracy

- Often significantly better than a single classifier derived from D
- For noisy data: not considerably worse, more robust
- Improved accuracy in prediction
- Decreases error by **decreasing the variance** in the results due to **unstable learners**
(some algorithms' output can change dramatically when the training data is slightly changed, e.g. decision trees, NNs, ...)
- **Increases classifier stability, reduces variance!**

(Breiman, 1996)

Random Subspace Method




Ensemble Learning – Overview

- Benefits of ensembles
- How to combine their outputs
- Bagging
- ▶ Boosting
 - AdaBoost

Ensemble Learning

- Another way of thinking about ensemble learning:
 - way of **enlarging the hypothesis space**, i.e., the ensemble itself is a hypothesis
 - the new hypothesis space is the set of all possible ensembles constructible from hypotheses of the original space
- Increased power of ensemble learning:



- Three linear threshold hypotheses (positive examples on the non-shaded sides)
- Ensemble classifies as positive any example classified positively by **all** three
- The **resulting triangular region** hypothesis is not expressible by any of the base hypotheses

Boosting

- How boosting works?

- $D_t(i)$ → • **Weights** are assigned to each training tuple i
- A series of k classifiers is iteratively learned ($t = 1, \dots, k$)
 - After a classifier M_t is learned, the weights of tuples are updated to allow the subsequent classifier, M_{t+1} , to **pay more attention to the training tuples that were misclassified** by M_t
 - The final M^* combines the votes of each individual classifier, where the **weight of each classifier's vote is a function of its accuracy**
- α_t →
- Boosting can be extended for numeric prediction
 - Compared with bagging:
Boosting tends to achieve greater accuracy,
but it also risks overfitting the model

Boosting: Strong And Weak Learners (1)

■ ***Strong Learner***

- Take labeled data for training
- Produce a classifier which can be ***arbitrarily accurate***
- Strong learners are an objective of machine learning

■ ***Weak Learner***

- Take labeled data for training
- Produce a classifier which is ***more accurate than random guessing***
- Weak learners can be base classifiers for ensemble methods

Boosting: Strong And Weak Learners (2)

- **Weak Learner:** only needs to generate a hypothesis with a training accuracy greater than 0.5, i.e., $< 50\%$ error over any distribution
 - Strong learners are very difficult to construct
 - Constructing weaker learners is relatively easy
- Can a set of **weak learners** create a single **strong learner**?
 - **Yes! Boost weak classifiers to a strong learner!** (Shapire, 1990)

Boosting: Use Weak Classifiers


Idea of iterative learning: Focus on difficult samples which are not correctly classified in the previous steps.

Use different data distribution:

- Start with **uniform weighting** of samples ($D_t(i)$)
- During each step of learning
 - **Increase weights** of the samples which are **not correctly** learned by the weak learner
 - **Decrease weights** of the samples which are **correctly** learned by the weak learner


Boosting Idea

**Weak
Classifier 1**




Boosting Idea

**Weights
Increased**




Boosting Idea




Boosting Idea

**Weights
Increased**




Boosting Idea



Boosting Idea

Final classifier is
a combination of
weak classifiers



Boosting: Combine Weak Classifiers

- Idea for combination: better weak classifier gets a larger weight!
- Weighted voting
 - Construct **strong classifier** by **weighted voting** of the weak classifiers
 - Weight of each learner is directly proportional to its accuracy

Ensemble Learning – Overview

- Benefits of ensembles
- How to combine their outputs
- Bagging
- Boosting


 AdaBoost

AdaBoost: Adaptive Boosting

- Does not need to know the number of weak classifiers in advance
- Does not need to know error bounds on the weak classifiers, unlike earlier boosting algorithms

AdaBoost: Adaptive Boosting

- Each rectangle corresponds to an example, with **weight proportional to its height**.
 $\swarrow D_t(i)$
- Crosses correspond to misclassified examples.
- Size of “decision tree” indicates **the weight of that hypothesis** in the final ensemble.
 $\swarrow \alpha_t$



Initialization

Given : $(x_1, y_1) \dots (x_n, y_n)$, where $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize distribution (weight) $D_{t=1}(i) = 1/n$; such that $n = M + L$

M = number of positive (+1) examples; L = number of negative (-1) examples

For $t = 1, \dots, T$

{ Step1a : Find the classifier $h_t : X \rightarrow \{-1, +1\}$ that minimizes the

error with respect to D_t , that means : $h_t = \arg \min_q [\varepsilon_q]$

Step1b : error $\varepsilon_t = \sum_{i=1}^n D_t(i) * I_{[h_t(x_i) \neq y_i]}$, where $I_{[h_t(x_i) \neq y_i]} = \begin{cases} 1 & \text{if } [h_t(x_i) \neq y_i] \text{ (classified incorrectly)} \\ 0 & \text{otherwise} \end{cases}$

checking step : prerequisite : $\varepsilon_t < 0.5$: (error smaller than 0.5 is ok) otherwise stop.

Step2 : $\alpha_t = \frac{1}{2} \ln \frac{1 - \varepsilon_t}{\varepsilon_t}$, α_t = weight (or confidence value).

Step3 : $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$, see next slide for explanation

Step4 : Current total cascaded classifier error $CE_t = \sum_{j=1}^t E_j(t, \alpha_j, h_j(x_i))$

while the current classifier error $E_t = \frac{1}{n} \sum_{i=1}^n I(t, \alpha_t, h_t(x_i))$,

and $I()$ is defined as follows :

If x_i is correctly classified by the current cascaded classifier , i.e.

$y_i = \text{sign} \left(\sum_{\tau=1}^t \alpha_\tau h_\tau(x_i) \right)$, hence error $I(t, \alpha_t, h_t(x_i)) = 0$

If x_i is incorrectly classified by the current cascaded classifier i.e.

$y_i \neq \text{sign} \left(\sum_{\tau=1}^t \alpha_\tau h_\tau(x_i) \right)$, hence error $I(t, \alpha_t, h_t(x_i)) = 1$

If $CE_t = 0$ then $T = t$, break;

}

The output $o_t(x_i) = \sum_{\tau=1}^t \alpha_\tau h_\tau(x_i)$, and $S(t, \alpha_t, h_t(x_i)) = \begin{cases} 1 & \text{if } y_i = \text{sign}(o_t(i)) \\ 0 & \text{otherwise} \end{cases}$

where Z_t = normalization factor, so D_t is a probability distribution

$$\begin{aligned} Z_t &= \sum_{i=1}^{n_{\text{correctly_classified}}} \text{correct_weight} + \sum_{i=1}^{n_{\text{incorrectly_classified}}} \text{incorrect_weight} \\ &= \sum_{i=1}^{n_{\text{correctly_classified}}} D_t(i) e^{-\alpha_t} y_i h_t(x_i) + \sum_{i=1}^{n_{\text{incorrectly_classified}}} D_t(i) e^{\alpha_t} y_i h_t(x_i) \end{aligned}$$

*enlarged
versions
on the
following
slides*

Main Loop

Final strong classifier: $H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$

Strong Classifier

Initialization

Given $(x_1, y_1), \dots, (x_n, y_n)$, where $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize weights of samples $D_{t=1}(i) = 1/n$;

such that $n = M + L$

M = number of positive (+1) examples;

L = number of negative (-1) examples

Adapted from:

Kin Hong Wong: Adaboost for building robust classifiers.

<http://appsrv.cse.cuhk.edu.hk/~khwong/>

Main Loop (Steps 1, 2, 3)

For $t = 1, \dots, T$

{

Step1a : Find the classifier $h_t : X \rightarrow \{-1, +1\}$ that minimizes the error with respect to D_t : $h_t = \arg \left[\min_q (\varepsilon_q) \right]$

Step1b : error $\varepsilon_t = \sum_{i=1}^n D_t(i) * I_{[h_t(x_i) \neq y_i]}$,

$I = \text{incorrectness}$

where $I_{[h_t(x_i) \neq y_i]} = \begin{cases} 1 & \text{if } [h_t(x_i) \neq y_i] \text{ (classified incorrectly)} \\ 0 & \text{otherwise} \end{cases}$

Check whether $\varepsilon_t < 0.5$, otherwise stop.

Step2 : $\alpha_t = \frac{1}{2} \ln \frac{1 - \varepsilon_t}{\varepsilon_t}$, $\alpha_t = \text{weight of classifier (confidence)}$.

Step3 : $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$

Main Loop (Step 4)

Step4 : Current total cascaded classifier error $CE_t = \sum_{j=1}^{j=t} E_j(t, \alpha_\tau, h_\tau(x_i))$

where the current classifier error $E_\tau = \frac{1}{n} \sum_{\tau=1}^n I(t, \alpha_\tau, h_\tau(x_i))$,

and $I()$ denotes incorrectness for x_i of the current cascaded classifier :

$$y_i = \text{sign}\left(\sum_{\tau=1}^t \alpha_\tau h_\tau(x_i)\right) \rightarrow I(t, \alpha_\tau, h_\tau(x_i)) = 0$$

$$y_i \neq \text{sign}\left(\sum_{\tau=1}^t \alpha_\tau h_\tau(x_i)\right) \rightarrow I(t, \alpha_\tau, h_\tau(x_i)) = 1$$

If $CE_t = 0$ then $T = t$, break;

add threshold
if needed

Final strong classifier: $H(x) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(x) - 0\right)$

Hybrid Ensemble Learning with the NAO



NAO learns objects based on an **ensemble** of **neural networks**

- Every network classifies based on **different features**:
pixel patterns, color & texture, or SURF features