

Data-driven Intelligent Systems

Lecture 17 Reinforcement Learning II



KNOWLEDGE
TECHNOLOGY

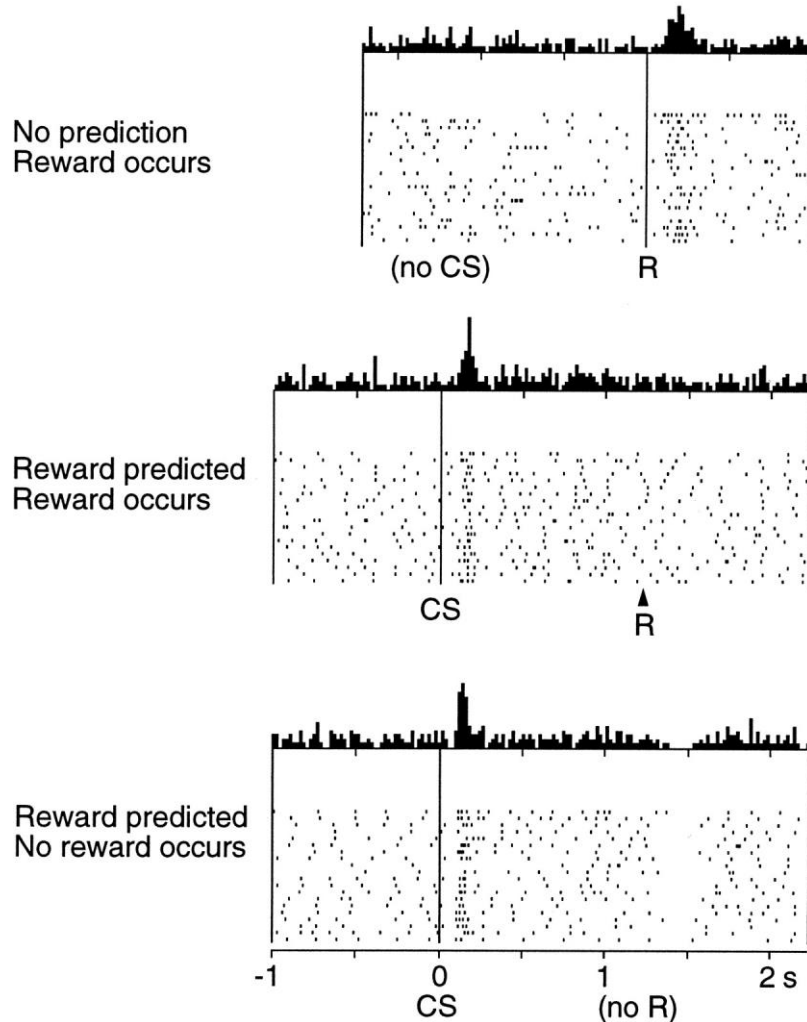
<http://www.informatik.uni-hamburg.de/WTM/>

Outline

- ▶ Dopamine in the Brain Relates to TD Error
 - Policy-gradient: REINFORCE
 - Goal-conditioned RL and Hindsight Experience Replay
 - Hierarchical RL
 - Model-based RL: MuZero

Biological Analogy: Dopamine signals TD error

Firing (spikes) of dopamine neurons



Pavlovian conditioning:

Initially: When a reward is given (e.g. juice) the animal reacts (e.g. salivates)

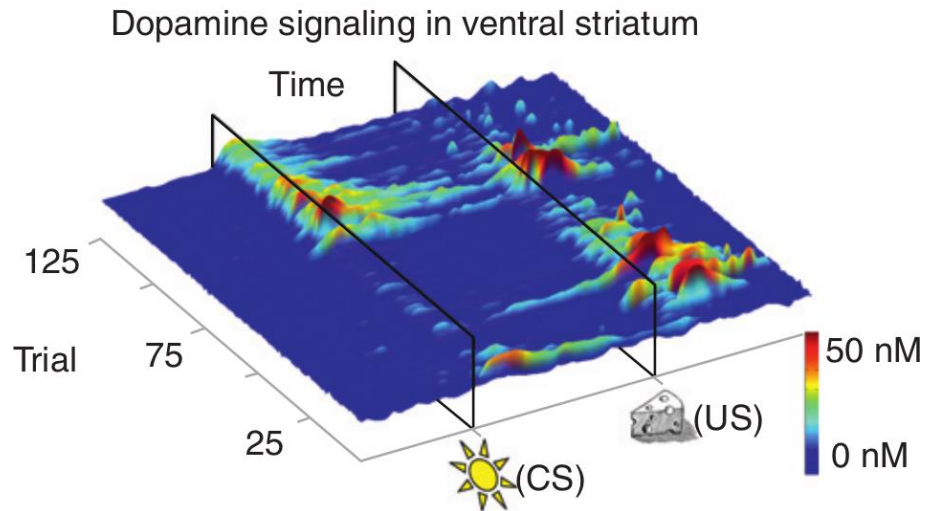
Learning: A neutral stimulus (e.g. bell) is consistently presented before giving the reward

Result: the animal reacts at the neutral stimulus, which becomes the conditioned stimulus (CS), but not during reward.

Interpretation: TD error is relevant.

Dopamine neurons behave in the same way.

Biological Analogy: Dopamine signals TD error



Further Neuro-transmitters

Good evidence:

- Dopamine suggested to correspond to the TD error

Rather speculative:

- Norepinephrine / Noradrenaline signals unexpected events, i.e. surprise or uncertainty

Dayan, Yu (2006) Phasic norepinephrine: A neural interrupt signal for unexpected events


- Acetylcholine signals unexpected uncertainty

Yu, Dayan (2003) Acetylcholine, Norepinephrine, and Spatial Attention

- Serotonine correlates with discount factor γ

Yoshida, Uchibe, Doya (2013) Reinforcement learning with state-dependent discount factor

Basal Ganglia as Brain Candidate for RL



Model hypothesis:

- Striatum encodes the state
- Thalamus receives the action selection
- Direct pathway is via inhibition of inhibition
- Dopamine learning signals arise from SN and VTA

SN = substantia nigra

VTA = ventral tegmental area

Reasons for an Agent to Act

- Rewards
 - Classical rewards, e.g. food & drink
 - Social & task rewards, e.g. praise, exam marks, salary
 - Negative rewards, e.g. pain
 - Epistemic rewards
 - Reward from information gain (~curiosity)
 - Reward for efficient sensory coding, e.g. better compression
- Imitation behavior
 - Inverse RL
- Homeostasis
 - Self-protection, e.g. reflexes
 - Daily rhythm with activity and rest/sleep

Outline

- Dopamine in the Brain Relates to TD Error

Policy-gradient: REINFORCE

- Goal-conditioned RL and Hindsight Experience Replay
- Hierarchical RL
- Model-based RL: MuZero

Policy Gradient

Previous lecture: Value-based RL

- Agent estimates $Q(s,a)$ or $V(s)$ value functions, and use these while also finding the best policy π

Alternative: Policy-gradient based RL

- Search best policy π directly by gradient ascent
- *“But this can’t work, because we have no learning signals during steps when the reward is zero? – We need a value function to solve this temporal credit assignment problem!”*
- Solution: Use wider estimates of Return; use NN or RNN (for interpolation and backprop. through time)
 - allows learning at timesteps other than when a reward is directly present

Policy Gradient - Equations

Objective is to maximize the expected future return R .

We take the gradient ∇ with respect to the parameters θ :

$$\nabla_{\theta} \langle R_{\pi_{\theta}} \rangle = \sum_s p(s) \sum_a \nabla_{\theta} \pi_{\theta}(a|s) \cdot Q(s, a)$$

$$\text{use: } \nabla_{\theta} \pi_{\theta}(a|s) = \pi_{\theta}(a|s) \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a|s)} = \pi_{\theta}(a|s) \nabla_{\theta} \text{Log } \pi_{\theta}(a|s)$$

$$\frac{d}{dx} \log x = \frac{1}{x}$$

$$= \underbrace{\sum_s p(s) \sum_a \pi_{\theta}(a|s)}_{\text{Expectation value, compute by sampling}} \cdot \underbrace{\nabla_{\theta} \text{Log } \pi_{\theta}(a|s)}_{\text{Compute by gradient back-propagation from output units}} \cdot \underbrace{Q(s, a)}_{\text{Use external rewards } R}$$

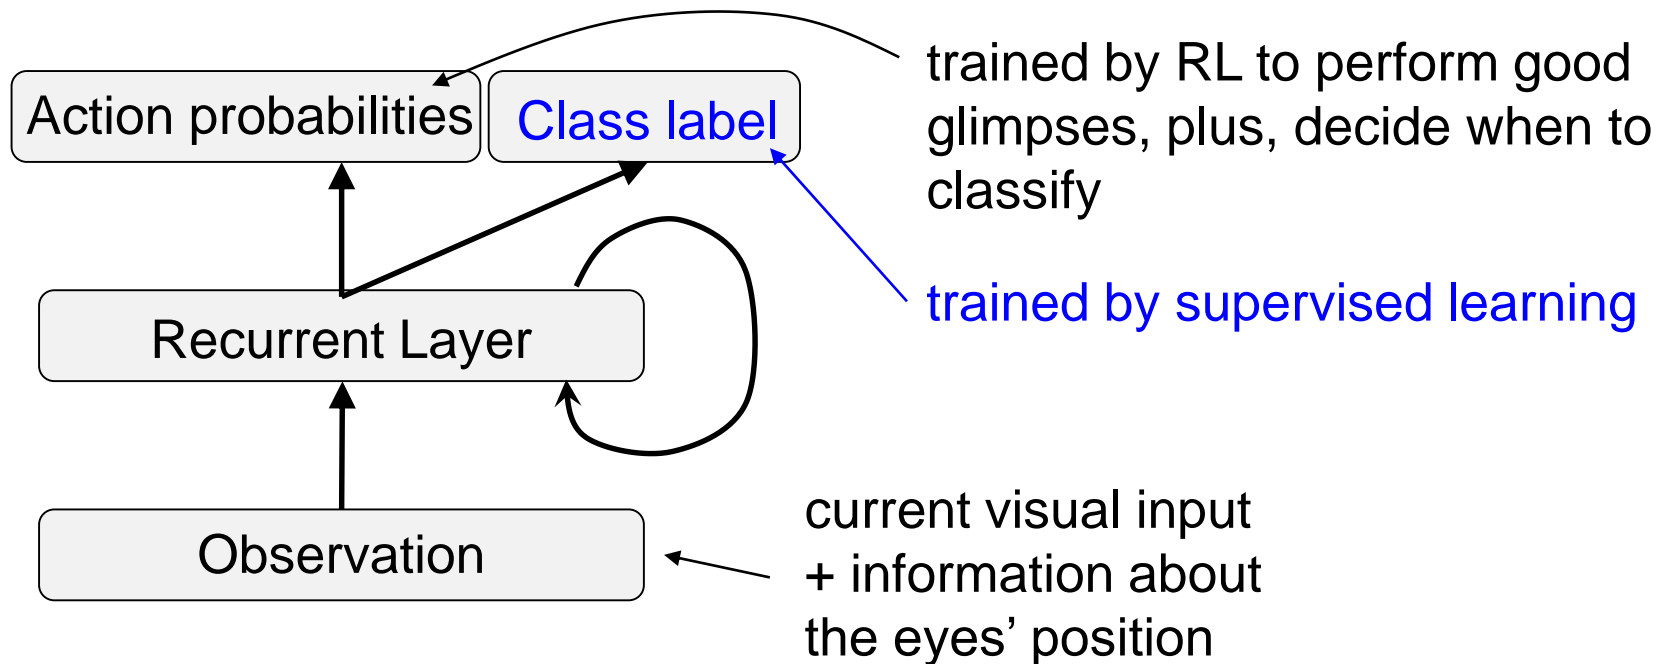

Variance reduction: subtract expected baseline value b : $R \leftarrow R - b$

Policy Gradient – Visual Attention

Task: visually identify a handwritten digit (MNIST)

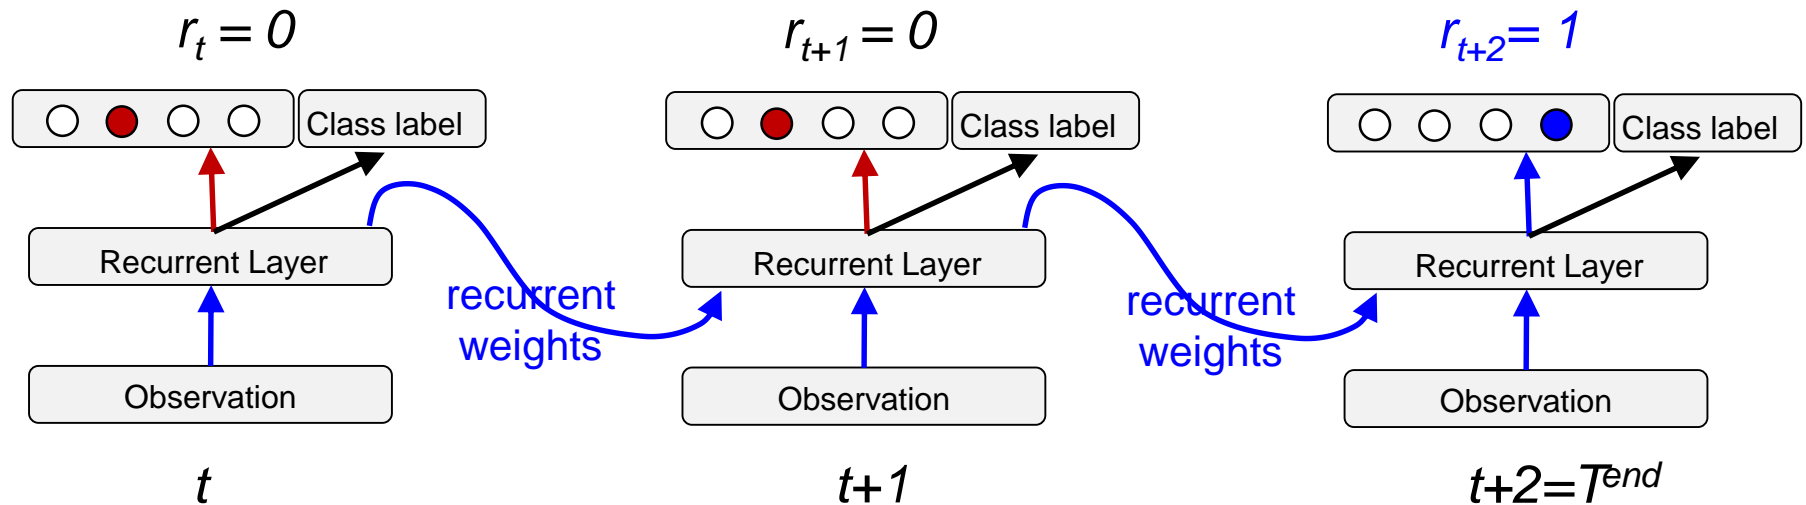

Constraint: good resolution only in small field of view

Strategy: learn a sequence to glimpse at certain positions; integrate information via RNN; then classify the digit



Policy Gradient – Visual Attention

Backpropagation through time pairs reward at T^{end} with past activations only in input and hidden layers



Train with $R = \sum r_t$ for entire sequence
→ weights to all used actions can be learnt

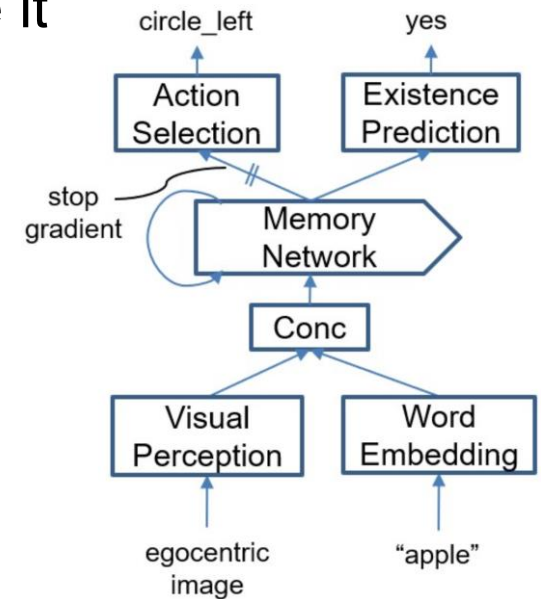

Related technique: eligibility traces

Policy Gradient – Occlusion Reasoning

Task: answer whether a specific object is on the table

Constraint: object might be occluded by another one

Learnt strategy: walk around the table, if the object is not visible, if it is small and if a larger object might occlude it



Policy Gradient – Occlusion Reasoning

Efficient Robotic Object Existence Prediction by Occlusion Reasoning

Supplementary Video

Mengdi Li, Cornelius Weber, Matthias Kerzel, Jae Hee Lee, Zheni Zeng, Zhiyuan Liu, Stefan Wermter

University of Hamburg and Tsinghua University



Outline

- Dopamine in the Brain Relates to TD Error
- Policy-gradient: REINFORCE
- ▶ Goal-conditioned RL and Hindsight Experience Replay
- Hierarchical RL
- Model-based RL: MuZero


Goal-conditioned RL

So far, problem of unflexibility:

- Typically, Q-values are assigned to states irrespective of goal
- Changing the goal requires new Q-values, hence new learning!
- Separate Q-values for every goal are costly, if a Q-table is being used

The solution is **goal-conditioned RL**:

- Enhance the input with additional units that encode the goal
 - Agent learns to reach various goals
 - Goals can be dynamically set
 - e.g. by another network
→ hierarchical RL



Input = (state, goal)

Hindsight Experience Replay (HER)

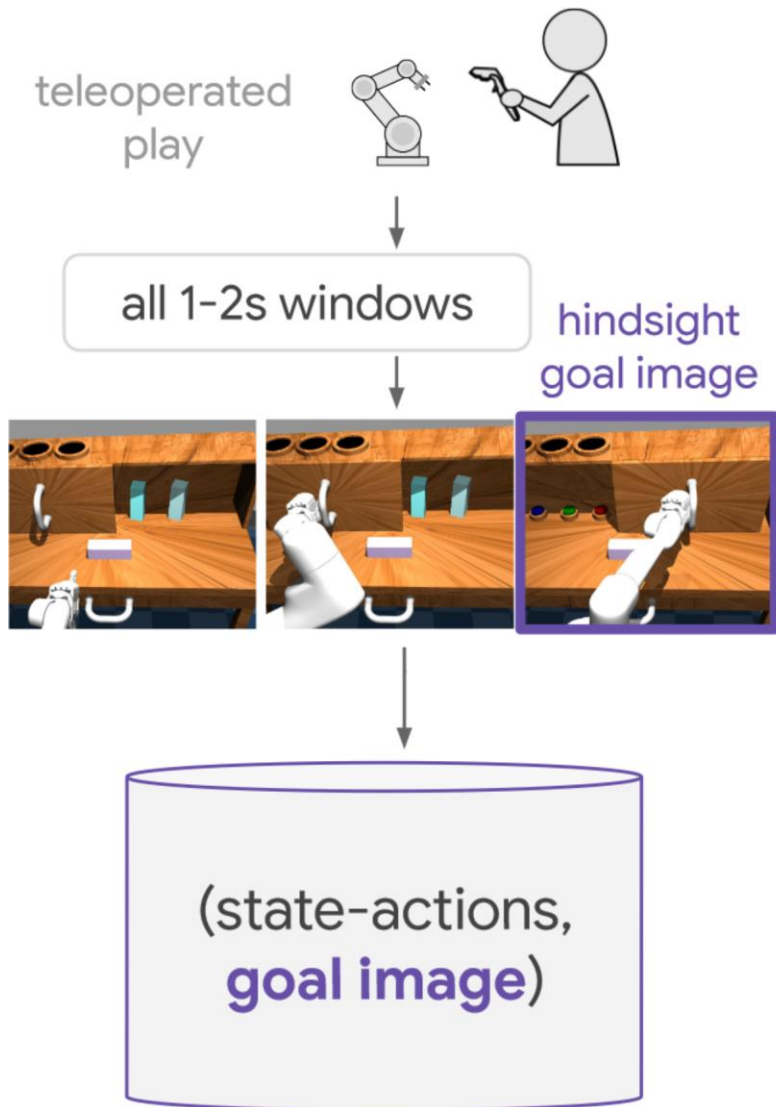
Recap: *Experience Replay*

- During environment exploration, memorize experiences
- Learn from randomly sampled experiences from memory

Hindsight Experience Replay

- Use architecture for goal-conditioned RL
 - Model input is state and goal: (s, g)
- Take any experience (typically, an extended sequence)
 - $(s_0, s_1, \dots, s_T) \leftarrow$ state sequence from time = 0 until T
- In hindsight, assume s_T had been the goal g
 - Model input: (s_0, s_T)
- This teaches the agent to reach s_T
- Will be repeated for many different $\{s_T\}$

HER Example



Data collection

- People remote-control the robot; playful behavior without particular instructions
- Nevertheless, people do useful things, such as opening doors

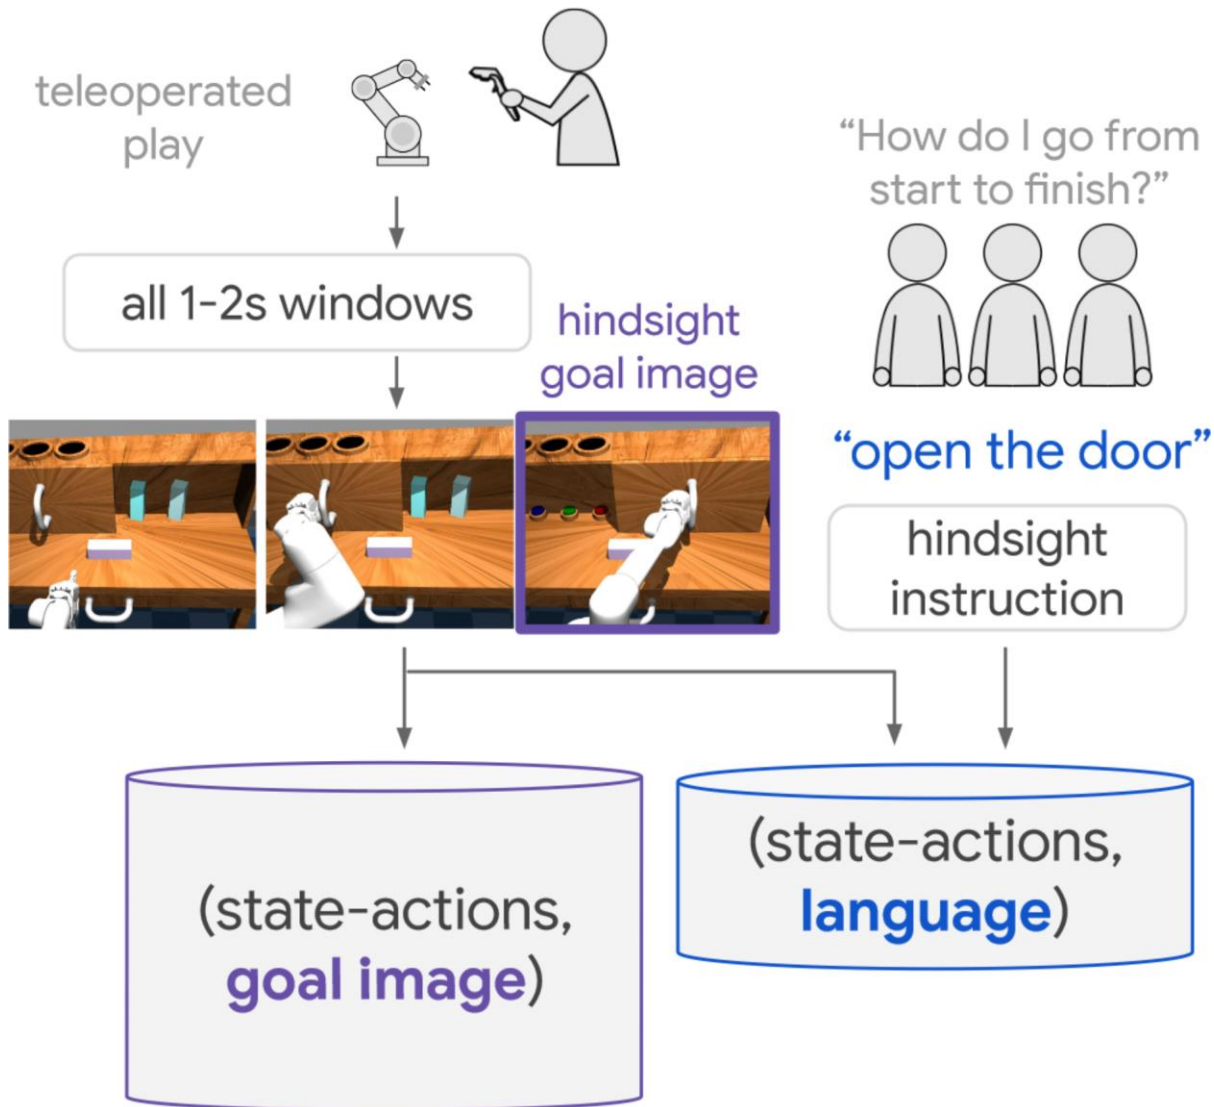
HER learning

- Short sequences presented
- Last frame is the goal
- Agent learns to reach any goal

Usage

- Need to show an image of the goal state; agent can then reach it

HER Example



Language extension

- Image goals are supplemented by language goals


HER learning

- Deep NN learns joint representation of image and language goal

Usage

- Lang. instruction provides the goal

HER Example



now: **do not do anything**

next:

Outline

- Dopamine in the Brain Relates to TD Error
- Policy-gradient: REINFORCE
- Goal-conditioned RL and Hindsight Experience Replay
- ▶ Hierarchical RL
 - Model-based RL: MuZero

Hierarchical RL

“Flat” RL has limitations with hierarchically structured tasks:

- i.e. consisting of subtasks
- particularly if the subtasks are not rewarded

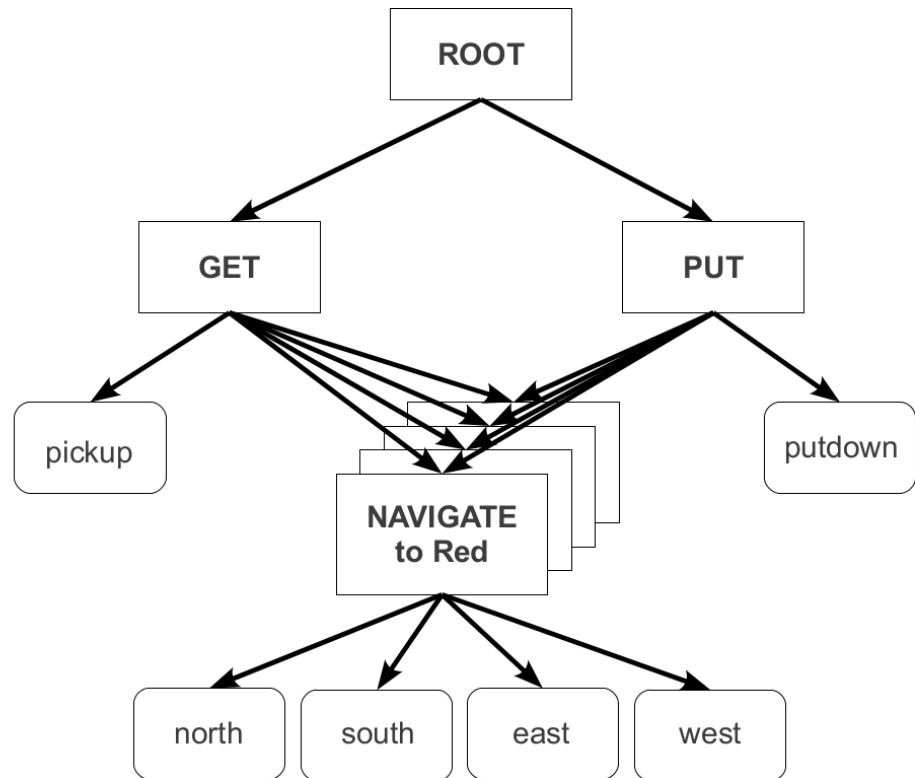

Examples in Atari games:

- must collect a key before a door can be opened
- multiple small missions to complete in a specific order

Hierarchical solution: decompose the task into subtasks:

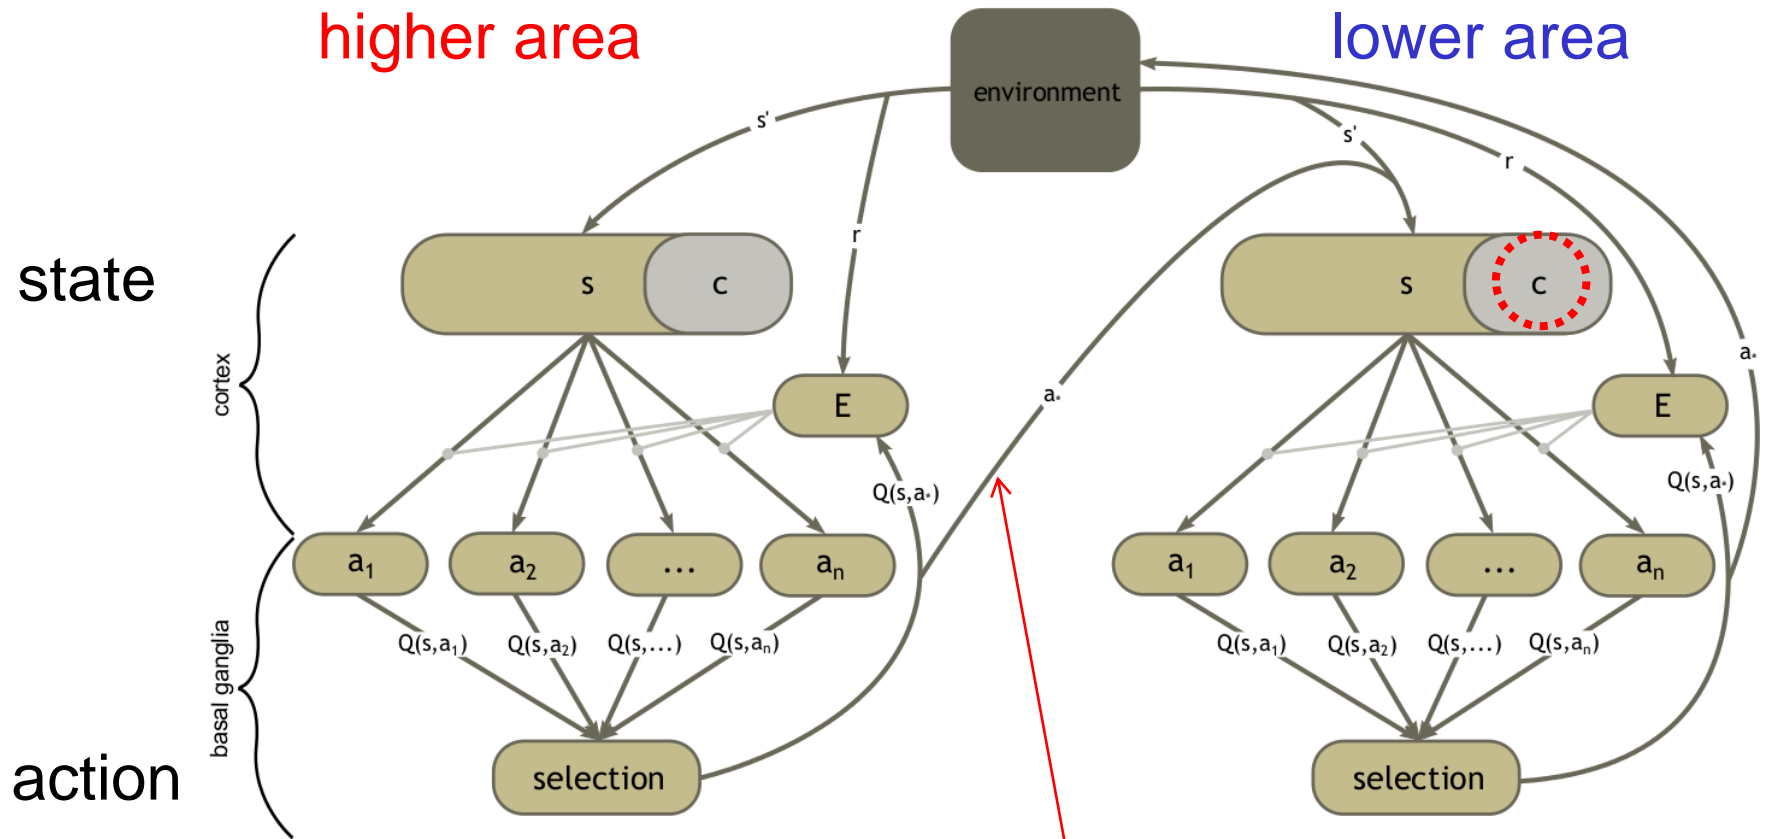
- on a low hierarchical level, solve the subtasks
- on a high hierarchical level, activate the subtasks and provide their respective goals

Hierarchical RL – taxi domain



- reward given when taxi puts down passenger at destination
- hierarchical decomposition makes the problem tractable
 - low level learns “navigate to X” as a subtask
 - high level learns to arrange subtasks and set their goals

Hierarchical RL – top-down control



top-down connections determine the current navigation goal (via context units c)

Hierarchical RL – top-down influence in biology





Figure: del Campo et al. (2013) A positron emission tomography study of nigro-striatal dopaminergic mechanisms underlying attention: implications for ADHD and its treatment. *Brain* 136 (11), 3252-3270.

Outline

- Dopamine in the Brain Relates to TD Error
 - Policy-gradient: REINFORCE
 - Goal-conditioned RL and Hindsight Experience Replay
 - Hierarchical RL
-  Model-based RL: MuZero


Model-based RL

- **Model** of the transition structure of the world

$$P(s_{t+1}) = P(s_{t+1}|s_t, a_t)$$

- for all states and all actions
- Can be learnt by exploration, independently of any rewards
- The model may also represent the reward structure $r(s_{t+1}, s_t)$
- *Use a model:*
 - compute all Q-values based on internal simulation by model
 - compute total Return for next selection, or an entire plan
 - compare different plans, to select the best one
- *Limitations:*
 - exponential number of states/actions; cycles
 - breadth-first vs. depth-first

Model-based RL




- tree search through future states and rewards

Model-based vs model-free RL

- What if the environment changes, e.g.
 - a reward is being devaluated
 - a transition becoming impossible
 - a new shortcut is introduced
- A **model-free** RL learner
 - will initially maintain its stimulus-response-like policy at all other states than the changed ones
 - must re-visit all states to update their values (iteratively)
- A **model-based** learner
 - can automatically adjust to changes in this information
- Humans used a mixture of model-free and model-based RL
 - habits vs planning

MuZero: Architecture




Three neural networks:

- Representation function h
 - maps raw observations to latent representations s_t
- Dynamics function g
 - maps (s_t, a_t) to s_{t+1} and predicts the reward r
- Policy and value prediction function f
 - akin actor-critic network

The dynamics function is a forward model (world model) and allows recursive prediction of future outcomes, i.e. planning; but: *in latent space*

MuZero: Planning by Monte-Carlo Tree Search



While exploring (in) the world, the agent plans ahead


- Use the dynamics function g to simulate future states s_{t+1}, s_{t+2}, \dots , of which f predicts their value
- Different actions a_t are explored
- MCTS extends various branches into an emerging planning tree; better branches will be more often chosen and extended
- Finally, **one** action from the root node is chosen to be executed, proportionally to the visit count of its branch

At next real time step: span another complete planning tree

MuZero: Training via Experience Replay

- All experiences are saved in a replay buffer:
 - sequences of (observations, actions, rewards)
- Learning happens from randomly chosen experiences
 - All learning-relevant values can be computed from current network parameters
 - Using forward- and backpropagation through the network's functions h , g , and f
- Planning happens only during “real world” exploration


MuZero: Architecture Extension



Four
~~Three~~ neural networks:

- Representation function h
 - maps raw observations to latent representations s_t
- Dynamics function g
 - maps (s_t, a_t) to s_{t+1} and predicts the reward r
- Policy and value prediction function f
 - akin actor-critic network
- Reconstruction function h^{-1}
 - Reconstructs the input. Allows to pretrain also h and g by unsupervised random exploration of the environment

MuZero: Extension for Continuous Actions




So far, discrete actions:

- Branching in MCTS happened along discrete action choices

Continuous action space:

- Exploration and branching happen by randomly selecting continuous action vectors
- These are chosen from a Gaussian probability density function, centered around the current most-likely action (according to the policy)
- The tree will still consist of discrete branches, as before

Continuous MuZero: Intercepting a Rolling Ball



Summary of RL

- TD Error, derived from theory, explains biological learning
- RL complements unsupervised & supervised learning
 - adds goal-directedness
- Dynamic programming solves temporal credit assignment problem
- Techniques of supervised learning are useful helpers:
 - MLP and deep architectures approximate return
- Hierarchical RL: slow high-level actions modulate/control fast low-level actions
- Model-based RL coexists with model-free RL